
Applied Data Analysis

R-Laboratory 5

Implementing Normal Linear Models – Linear Models Beyond Normality Analysis of Variance

Useful packages and functions:

- | | | |
|-------------------------------|--------------------------|----------------------------------|
| • <code>model.matrix()</code> | • <code>anova()</code> | • <code>car::leveneTest()</code> |
| • <code>qr()</code> | • <code>rexp()</code> | • <code>aov()</code> |
| • <code>qr.solve()</code> | • <code>density()</code> | • <code>TukeyHSD()</code> |
| • <code>chol2inv()</code> | • <code>anova()</code> | • <code>par()</code> |
| • <code>qt()</code> | • <code>car</code> | |

Task 18

- Load the `.RData` file of the pre-processed data of *Solar* (Task 10) into the R workspace.
- Create four boxplots for the attribute `Pmax`, one for each batch.
- Carry out an analysis of variance for the attribute `Pmax` regarding the factor `batch`.
- Analyze the model assumptions as in Task 14 (c). Additionally, test on level $\alpha = 0.05$, if there is evidence against the assumption of equal variance in each group of factor level.

Hint: Use a Levene-Test with the R function `leveneTest` from the package `car` to test the equality of the variances.

- Test on the significance level $\alpha = 0.05$, if the null hypothesis of equal means of `Pmax` for the four batches is rejected.

Hint: You may call the function `summary` with the fitted model as argument.

- Carry out a pairwise comparison of the batches regarding `Pmax` using a Tukey-Test on significance level $\alpha = 0.1$. Create a plot of the computed confidence intervals.

Task 19

- Consider the data set `ToothGrowth` from the package `datasets` and transform the attribute `dose` to type `factor`.
- Create a boxplot for the value of `len`, separated on all factor combinations of `supp` and `dose`. What is your impression?

- (c) Fit a linear model for `len`, where `supp` and `dose` (inclusive interaction) are explanatory variables.
- (d) Analyze the model assumptions as in Task 14 (c). Additionally, test on level $\alpha = 0.05$, if there is evidence against the assumption of equal variance in each group of factor combination.
- (e) If the model assumptions are sufficiently satisfied, test on level $\alpha = 0.05$, if there is any influence of the explanatory variables on the value of `len`.
- (f) If there is any influence of explanatory variables, analyze it concretely: Test on overall $\alpha = 0.05$, if the interaction of `supp` and `dose` and possibly their main effect has an influence on the value of `len`.
Hint: Use a Bonferroni correction to ensure that all tests together satisfy the significance level α . This means that you can compare each p-value with an adapted significance level $\tilde{\alpha} = \frac{\alpha}{k}$, where $k \in \mathbb{N}$ is the number of tests.
- (g) If some effects are not significantly different from zero, fit a new model with the removed terms to get the final model for interpretation.

Task 20

Let $\mathbf{y} = (y_1, \dots, y_n)'$ be a realization of the random sample $\mathbf{Y} = (Y_1, \dots, Y_n)'$. For \mathbf{y} , consider the linear model $\mathbf{Y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\varepsilon}$, where \mathbf{X} is a fixed $(n \times d)$ -model (or design) matrix with $d < n$ and rank d , $\boldsymbol{\beta}$ is a $(d \times 1)$ -vector of model parameters and $\boldsymbol{\varepsilon} \sim \mathcal{N}_n(\mathbf{0}, \sigma^2 I_n)$ with $\sigma^2 > 0$.

Implement your own linear model R-function. This function should use a model formula and a `data.frame` object as input and should deliver the following output: least squares estimates $\hat{\boldsymbol{\beta}}$ of model parameters $\boldsymbol{\beta}$, $\text{var}(\hat{\boldsymbol{\beta}})$, the unbiased estimator $\hat{\sigma}^2$ of σ^2 , and $R^2 = \frac{SSR}{SST}$.

Remark: Here we write $\|\mathbf{y} - y^*\|_2^2 = \|\hat{\boldsymbol{\mu}} - y^*\|_2^2 + \|\mathbf{y} - \hat{\boldsymbol{\mu}}\|_2^2$ as $SST = SSR + SSE$ (total sum of squares = sum of squares due to the regression + sum of squared errors), where $y^* = \bar{y}$ for a model with intercept parameter and $y^* = 0$ for models without.

Hints:

- (a) Use the function `model.matrix` to create the design matrix from the right hand side of the given formula.
- (b) Let $\mathbf{X} = \mathbf{QR}$ be a \mathbf{QR} -decomposition of the matrix \mathbf{X} into a upper triangular $(d \times d)$ -matrix \mathbf{R} and a matrix \mathbf{Q} of the first d columns of an orthogonal $(n \times n)$ -matrix. Use the fact that the least squares estimator is given by $\hat{\boldsymbol{\beta}} = \mathbf{R}^{-1}\mathbf{Q}'\mathbf{y}$ to implement a numerically more stable procedure than that using the classical representation $\hat{\boldsymbol{\beta}} = (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{y}$.

Test your function on the data set `Solar` and calculate 90%-confidence intervals for the differences $\alpha_{i+1} - \alpha_i$, $i = 1, \dots, 3$, of the factor levels $1, \dots, 4$ of `batch` treated as factor (cf. Task 15).

Task 21

Set the seed to 2020 and repeat the following procedure $n = 1000$ times for $N = 10, 20, 50$:

- (a) Generate samples x_{11}, \dots, x_{1N} from a uniform distribution on $(0, 40)$.
- (b) Generate samples x_{21}, \dots, x_{2N} from $\mathcal{N}(15, 10^2)$.
- (c) Generate error values $\varepsilon_1, \dots, \varepsilon_N$ from the distribution of the random variable ε , where

$$\frac{\varepsilon}{5} + 1 \sim \text{Exp}(1)$$

and $\text{Exp}(1)$ denotes the exponential distribution with parameter 1.

- (d) Generate a sample y_1, \dots, y_N by setting

$$\mu_i = \beta_1 + \beta_2 x_{1i} + \beta_3 x_{2i}, \quad i = 1, \dots, N,$$

where $\beta_1 = 35$, $\beta_2 = 0.5$, $\beta_3 = -0.1$ and $y_i = \mu_i + \varepsilon_i$.

- (e) Estimate $\boldsymbol{\beta} = (\beta_1, \dots, \beta_3)$ using the least squares estimator $\hat{\boldsymbol{\beta}}$.
- (f) Compute the norm-difference $\|\boldsymbol{\beta} - \hat{\boldsymbol{\beta}}\|$ of the true parameter vector and its estimate.

Store the values of $\hat{\beta}_2$ and $\|\boldsymbol{\beta} - \hat{\boldsymbol{\beta}}\|$ for all n generated data sets and all $N = 10, 20, 50$. Then illustrate some results using the following graphics.

- (i) Create boxplots of the $\|\boldsymbol{\beta} - \hat{\boldsymbol{\beta}}\|$ for the three values of N .
- (ii) Create plots of the estimated densities of $\hat{\beta}_2$ for the three values of N and add the curve of the $\mathcal{N}(\hat{\mu}, \hat{\sigma}^2)$ density, where $\hat{\mu}, \hat{\sigma}^2$ are the mean value and the standard deviation of the sample of $\hat{\beta}_2$ values for the current N .

What do you observe?

```
#####  
#####TASK15#####  
#####
```

```
timeW.fit = lm(timeW ~ climb + distance, data = race_task15)
```

```
#b)
```

```
timeW.fit2 = lm(timeW ~ distance, data = race_task15)
```

```
#
```

```
# Analysis of Variance Table
```

```
#
```

```
# Model 1: timeW ~ distance
```

```
# Model 2: timeW ~ climb + distance
```

```
# Res.Df  RSS Df Sum of Sq  F  Pr(>F)
```

```
# 1    66 30686
```

```
# 2    65 12675  1    18011 92.36 4.223e-14 ***
```

```
# ---
```

```
# Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
# reject the null hypothesis that climb is not significant
```

```
distance = race_task15$distance
```

```
climb = race_task15$climb
```

```
ones_task15 = rep(1, nrow(race_task15))
```

```
mens_dat = cbind(ones_task15, climb, distance)
```

```
predicted_timeM = mens_dat %*% timeW.fit$coefficients
```

```
# which(race_task15["timeW"] == 490.05)
```

```
# [1] 41
```

```
# > predicted_timeM[41]
```

```
# [1] 456.1487
```

```
#
```

```
# #d)
```

```
# > cor(race_task15$timeM, race_task15$timeW)
```

```
# [1] 0.9958732
```

```
cor(Races$timeW, Races$timeM)
```

```
# we have strong positive correlation
```

```
# this indicates that in a race where the males need more time, the females will also need more time
```

```
#e)
```

```
fit.timeMW = lm(timeM ~ -1 + timeW, data = race_task15)
```

```
#####
#####TASK18#####
#####
solar_task18 = read.table("R-Lab-Datasets/solar_task10.csv", header = TRUE,
sep = ",")

solar_task18$batch = as.factor(solar_task18$batch)

#b) boxplot for Pmax for each batch
ggplot(solar_task18, aes(x = batch, y = Pmax, fill=(batch), group = batch)) +
geom_boxplot(outlier.colour="red", outlier.shape=8,outlier.size=4)
#or this
#solar_boxplot = boxplot(solar_task18$Pmax ~ solar_task18$batch)

#anova of the Pmax corresponding for each batch
solar_aov = aov(solar_task18$Pmax ~ solar_task18$batch)

#levene test
levene.test.solar = leveneTest(solar_task18$Pmax ~ solar_task18$batch)

#plot the 4
# fitted vs residuals
# fitted vs standard residuals
# normal QQ
# residual vs leverage
par(mfrow = c(2,2))

plot(solar_aov)

par(mfrow = c(1,1))

#shapiro test
shapiro.test(residuals(solar_aov))

#Tukey test and its plot
plot(TukeyHSD(solar_aov, conf.level=0.9))
```

```
#####  
#####TASK19#####  
#####
```

```
#a) load toothgrowth data  
tooth.data = ToothGrowth
```

```
#to factorize the dose  
tooth.data$dose = as.factor(tooth.data$dose)
```

```
#b)  
par(mfrow = c(1,1))  
box_tooth_len_dose = boxplot(tooth.data$len ~ tooth.data$dose *  
tooth.data$supp)
```

```
#c) fit a model for len based on supp and dose
```

```
len.fit = lm(len ~ dose * supp, data = tooth.data)
```

```
#d)  
par(mfrow = c(2,2))  
plot(len.fit)
```

```
par(mfrow = c(1,1))  
aov.len = aov(len.fit)
```

```
#f) pairwise t test and fit another model by
```

```
# excluding less effective parameter
```

```
new.len.fit = lm(len ~ dose + supp, data = tooth.data)
```

```
pairwise.t.test(tooth.data$len, tooth.data$dose, p.adjust.method = "bonferroni")
```


#####TASK21#####
#####

set.seed(2020)

#a)
unif.sample = runif(50, 0, 40)
#b)
norm.sample = rnorm(50, 15, 10)
#c)
exp.sample = (rexp(50, 1)-1)*5
#d)
beta1 = 35
beta2 = 0.5
beta3 = -0.1

betas = c(beta1, beta2, beta3)

mu = beta1 + unif.sample*beta2 + norm.sample*beta3
Y = mu + exp.sample

#e)
pseudo_fit = lm(Y ~ unif.sample + norm.sample)

beta.hat = pseudo_fit\$coefficients

#f)
diff = sum((beta.hat - betas)^2)

```
#####  
#####TASK20#####  
#####
```

```
my.lm = function(formula, data){
```

```
Y = data[,as.character(formula[[2]])]#achieve the target column
```

```
X = model.matrix(formula, data = data)#generate the design matrix
```

```
parameter.name = colnames(data)
```

```
model.with.intercep=any(colnames(X)=="(Intercept)")
```

```
X=matrix(X,ncol=ncol(X))
```

```
p.columns=ncol(X)
```

```
n.observations=nrow(X)
```

```
degrees.of.freedom=n.observations-p.columns
```

```
qr.decomp.X=qr(X)
```

```
beta.hat=qr.solve(qr.decomp.X,y)
```

```
beta.hat=matrix(beta.hat)
```

```
mu.hat=X%*%beta.hat
```

```
residuals=y-mu.hat
```

```
SSE=sum((residuals)**2)
```

```
y.star=ifelse(model.with.intercep,mean(y),0)
```

```
TSS=sum((y-y.star)**2)
```

```
sigma.hat.2=SSE/degrees.of.freedom
```

```
R.2=(TSS-SSE)/TSS
```

```
beta.hat.cov=sigma.hat.2 * chol2inv(qr.decomp.X$qr)
```

```
rownames(beta.hat)=param.names
```

```
rownames(beta.hat.cov)=param.names
```

```
colnames(beta.hat.cov)=param.names
```

```
erg=list(beta.hat=beta.hat,beta.hat.cov=beta.hat.cov,R.2=R.2,sigma.hat.2=sigm  
hat.2)
```

```
}
```